GC Thrashing and Memory Pressure in a JVM

How to detect **GC Thrashing** and **Memory Pressure** in a JVM, covering:

- **Q** GC internals
- Allocation & promotion diagnostics
- Impact on CPU, latency, and throughput
- Open jstat interpretation
- Jacob Thread dump correlation
- MAT + GC log correlation

1. What is GC Thrashing — Technically

GC thrashing happens when the JVM cannot keep up with the allocation rate, leading to:

- Old Generation fills up rapidly, but GC fails to reclaim enough space
- High GC frequency and duration with little pause between events
- Application threads spend most of their time waiting at GC safepoints

Root Cause Mechanics:

- High object allocation rate (e.g., new objects per second)
- High **object promotion rate** due to premature tenuring
- Survivor spaces too small → forced promotions to Old Gen
- Fragmentation in Old Gen → insufficient contiguous memory → Full GC
- GC unable to compact or reclaim → JVM ends up in a near-continuous GC cycle

2. Memory Pressure — Deep Dive

Memory pressure arises when:

The allocated heap (-Xmx) is not enough for current workload

- Objects live longer than expected → Tenured fills up fast
- GC cannot free up enough space → Allocation failure
- Safepoint stalls block app threads during GC coordination

Indicators:

Symptom	Technical Description
FGC every 5–10 seconds	Old Gen is exhausted, not reclaimable
YGC with high Eden %	Eden fills rapidly – high allocation rate
Promotion Failure	Survivors or Old Gen can't accommodate
CMS concurrent mode failure	CMS couldn't finish GC before tenured filled up
G1 to-space exhausted	G1 couldn't reserve contiguous regions
Evacuation failure	G1GC failed during object copying

3. GC Log Analysis — Advanced Interpretation

Sample Log (G1GC):

2024-04-19T19:03:23.456+0000: 845.123: [GC pause (G1 Evacuation Pause) (young), 0.125s]

[Parallel Time: 101.0 ms, GC Workers: 8]

[GC Worker Start (ms): Min: 845123.4, Avg: 845123.6, Max: 845123.8, Diff: 0.4]

[Eden: 384.0M(384.0M)->0.0B(320.0M) Survivors: 40.0M->48.0M Heap: 1024.0M(2048.0M)->752.0M(2048.0M)]

What to Notice:

- Frequent young GC → Eden fills fast (384 MB to 0 MB)
- Survivor space is small (40M \rightarrow 48M) \rightarrow Possible **promotion pressure**
- Heap occupancy stays high → Old Gen not shrinking

4. jstat Analysis — Real-Time Deep Inspection

Run:

jstat -gcutil <pid> 1s 100

Sample output:

SO S1 E 0 M CCS YGC YGCT FGC FGCT GCT

5.00 60.00 85.00 92.00 97.00 95.00 309 12.34 85 45.56 57.90

Advanced Readings:

Metric	Diagnosis
E > 80% continuously	Eden is too small or allocation rate is high
O > 85% steadily	Old Gen saturation → Full GC trigger
YGC increases rapidly	High churn, frequent Young GCs
FGC count increases rapidly	GC thrashing; can't reclaim OldGen
FGCT vs GCT → >80% of time	JVM stuck in GC

Pattern to confirm GC thrashing:

- FGC increasing every few seconds
- O column stays high
- GCT climbing steeply

5. Heap Dump + GC Log Correlation

1. Generate dump during thrashing:

jmap -dump:live,format=b,file=thrashing-heap.hprof <pid>

- 2. Use Eclipse MAT:
- Load HPROF
- Open Dominator Tree
- Sort by Retained Heap
- Look for:

- HashMap, ThreadLocal, ArrayList with high retained size
- Static references holding memory
- Duplicate classloaders (in redeploy scenarios)

Example:

Class: com.example.cache.CustomCache \rightarrow Retained Heap: 512MB

► Red flag: cache is not releasing memory → memory pressure

Now correlate the timestamp in GC logs with heap dump snapshot.

ii 6. APM and JMX Deep Monitoring Strategy

Metrics to Alert On:

Metric	Threshold
Heap Used (%)	>85% consistently
GC Time (%)	>60-70%
FGC Count	>3 in 5 min
GC Pause Time	>1s frequently
Allocation Rate	>100MB/s
Promotion Rate	>50MB/s
Tenured Generation Utilization	>80% consistently
Survivor Overflow	Frequent
Eden Reclaim %	Low (<30%)

Tools:

- Prometheus + Grafana Dashboards
- Dynatrace memory hotspots
- New Relic GC metrics
- JConsole + VisualVM

7. Thread Dump Correlation

Run:

jstack -l <pid> > tdump.log

Check:

- **Application threads blocked on GC safepoint:**
- "http-nio-8080-exec-2" #34 waiting on condition
- Լ sun.misc.Unsafe.park()
- L GCLocker::stall

This means GC is freezing threads too often \rightarrow throughput loss.

- GC threads are highly active:
- "GC Thread#2" daemon prio=10 tid=0x00007f89dc01b800 runnable

If GC threads dominate CPU and app threads are mostly waiting → GC thrashing confirmed.



♦ 8. Real-Time JVM Allocation Profiling (Bonus)

Use Java Flight Recorder (JFR) or VisualVM with allocation sampling:

- Profile which classes are allocating most objects
- Check:
 - byte[] → buffer/caching issues
 - String → unbounded request data
 - List → holding references too long



9. JVM GC Tuning Strategy (if thrashing detected)

GC	Fix
CMS	Tune -XX:CMSInitiatingOccupancyFraction=40

G1	Tune -XX:InitiatingHeapOccupancyPercent=30
G1	Use -XX:MaxGCPauseMillis=200 to trigger earlier
G1	Avoid -XX:+UseLargePages unless memory pinned
General	Increase -Xms and -Xmx to avoid heap resizing

★ 10. Summary: How to Detect and Confirm GC Thrashing

Layer	Check
GC Logs	Frequent full GCs, low memory reclaimed
jstat	FGC/FGCT and O% trends
APM	High GC % time, pause durations
Heap Dump	Memory leaks or high-retention objects
Thread Dump	Threads waiting on GC, GC threads dominant
JFR	Allocation profiling shows high churn
Logs	java.lang.OutOfMemoryError, GC overhead limit exceeded